

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df = pd.read_csv("Housing.csv")
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   price                545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	price	area	bedrooms	bathrooms	stories	parking
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

```
In [6]: df.shape
```

```
Out[6]: (545, 13)
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: price                0
area                        0
bedrooms                   0
bathrooms                  0
stories                    0
mainroad                   0
guestroom                  0
basement                   0
hotwaterheating            0
airconditioning            0
parking                    0
prefarea                   0
furnishingstatus           0
dtype: int64
```

```
In [8]: a = df.select_dtypes(include=["float64", "int64"])
a.corr()
```

Out[8]:

	price	area	bedrooms	bathrooms	stories	parking
price	1.000000	0.535997	0.366494	0.517545	0.420712	0.384394
area	0.535997	1.000000	0.151858	0.193820	0.083996	0.352980
bedrooms	0.366494	0.151858	1.000000	0.373930	0.408564	0.139270
bathrooms	0.517545	0.193820	0.373930	1.000000	0.326165	0.177496
stories	0.420712	0.083996	0.408564	0.326165	1.000000	0.045547
parking	0.384394	0.352980	0.139270	0.177496	0.045547	1.000000

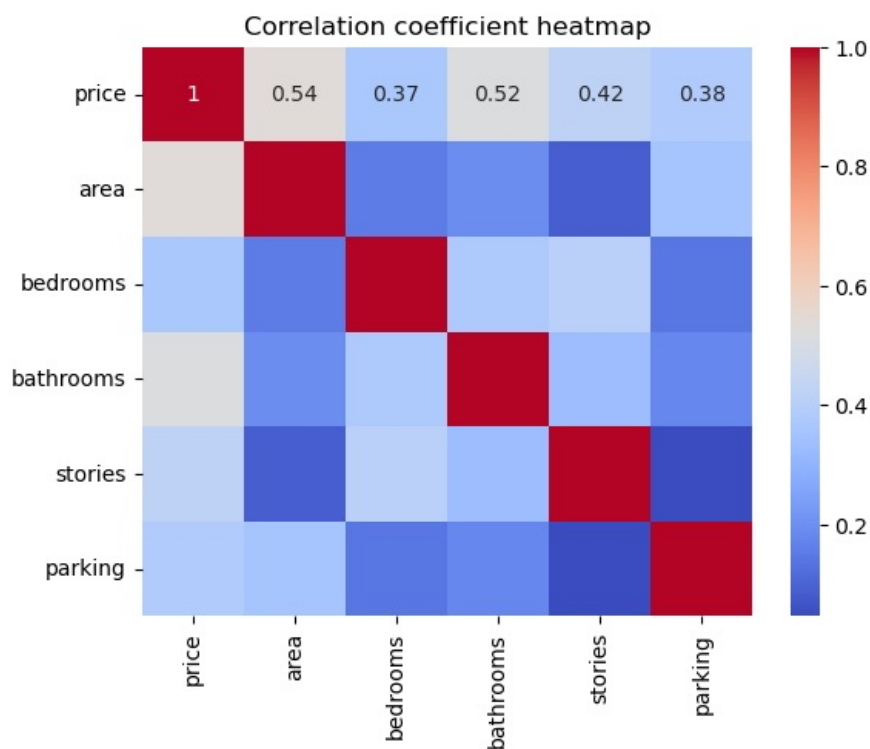
```
In [9]: c = df.select_dtypes(include=["object"])
c
```

Out[9]:

	mainroad	guestroom	basement	hotwaterheating	airconditioning	prefarea	furnishingstatus
0	yes	no	no	no	yes	yes	furnished
1	yes	no	no	no	yes	no	furnished
2	yes	no	yes	no	no	yes	semi-furnished
3	yes	no	yes	no	yes	yes	furnished
4	yes	yes	yes	no	yes	no	furnished
...	...	...	...	...	...	...	...
540	yes	no	yes	no	no	no	unfurnished
541	no	no	no	no	no	no	semi-furnished
542	yes	no	no	no	no	no	unfurnished
543	no	no	no	no	no	no	furnished
544	yes	no	no	no	no	no	unfurnished

545 rows × 7 columns

```
In [10]: sns.heatmap(a.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation coefficient heatmap")
plt.show()
```



## Linear regression models

```
In [11]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from scipy import stats
lr = LinearRegression()
price = df[["price"]]
for i in range(5):
    b = df[a.columns[i + 1]]
```

```

lr.fit(price, b)
predict1 = lr.predict(price)
r2 = r2_score(b, predict1)
corr, p_value = stats.pearsonr(df["price"], b)
print(f"Column: {a.columns[i + 1]}")
print(f"Correlation Coefficient: {corr:.2f}")
print(f"P-value: {p_value:.5f}")
print(f"R-squared Score: {r2:.2f}")
print("-----")

```

Column: area  
Correlation Coefficient: 0.54  
P-value: 0.00000  
R-squared Score: 0.29  
-----

Column: bedrooms  
Correlation Coefficient: 0.37  
P-value: 0.00000  
R-squared Score: 0.13  
-----

Column: bathrooms  
Correlation Coefficient: 0.52  
P-value: 0.00000  
R-squared Score: 0.27  
-----

Column: stories  
Correlation Coefficient: 0.42  
P-value: 0.00000  
R-squared Score: 0.18  
-----

Column: parking  
Correlation Coefficient: 0.38  
P-value: 0.00000  
R-squared Score: 0.15  
-----

```


In [12]: a = pd.get_dummies(df, columns=["mainroad", "guestroom", "basement", "hotwaterheating", "airconditioning", "prefarea_y
b = ["mainroad_yes", "guestroom_yes", "basement_yes", "hotwaterheating_yes", "airconditioning_yes", "prefarea_y
a

```

Out[12]:

	price	area	bedrooms	bathrooms	stories	parking	mainroad_no	mainroad_yes	guestroom_no	guestroom_yes	...	bas
0	13300000	7420	4	2	3	2	False	True	True	False	...	
1	12250000	8960	4	4	4	3	False	True	True	False	...	
2	12250000	9960	3	2	2	2	False	True	True	False	...	
3	12215000	7500	4	2	2	3	False	True	True	False	...	
4	11410000	7420	4	1	2	2	False	True	False	True	...	
...	...	...	...	...	...	...	...	...	...	...	...	...
540	1820000	3000	2	1	1	2	False	True	True	False	...	
541	1767150	2400	3	1	1	0	True	False	True	False	...	
542	1750000	3620	2	1	1	0	False	True	True	False	...	
543	1750000	2910	3	1	1	0	True	False	True	False	...	
544	1750000	3850	3	1	2	0	False	True	True	False	...	

545 rows × 21 columns



```

In [15]: for i in b:
lr.fit(price, a[b])
predict1 = lr.predict(price)
corr = df['price'].corr(a[i])
r = corr**2
print(f"Column: {i}")
print(f"Correlation Coefficient: {corr}")
print(f"R-squared Score: {r:.2f}")
print("-----")

```

Column: mainroad\_yes  
Correlation Coefficient: 0.2968984892639765  
R-squared Score: 0.09  
-----  
Column: guestroom\_yes  
Correlation Coefficient: 0.25551728993499967  
R-squared Score: 0.07  
-----  
Column: basement\_yes  
Correlation Coefficient: 0.18705659793805252  
R-squared Score: 0.03  
-----  
Column: hotwaterheating\_yes  
Correlation Coefficient: 0.09307284392139682  
R-squared Score: 0.01  
-----  
Column: airconditioning\_yes  
Correlation Coefficient: 0.4529540842560473  
R-squared Score: 0.21  
-----  
Column: prefarea\_yes  
Correlation Coefficient: 0.32977704986810746  
R-squared Score: 0.11  
-----

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js